

Import the dataset and do usual data analysis steps like checking the structure & characteristics of the dataset.

```
In [33]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from scipy.stats import norm
```

```
In [29]: df = pd.read_excel(r"C:\Users\Hp\Downloads\walmart.xlsx")
```

```
In [35]: df.head(20)
```

```
Out[35]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years
--	---------	------------	--------	-----	------------	---------------	----------------------------

0	1000001	P00069042	F	0-17	10	A	;
1	1000001	P00248942	F	0-17	10	A	;
2	1000001	P00087842	F	0-17	10	A	;
3	1000001	P00085442	F	0-17	10	A	;
4	1000002	P00285442	M	55+	16	C	;
5	1000003	P00193542	M	26-35	15	A	;
6	1000004	P00184942	M	46-50	7	B	;
7	1000004	P00346142	M	46-50	7	B	;
8	1000004	P0097242	M	46-50	7	B	;
9	1000005	P00274942	M	26-35	20	A	;
10	1000005	P00251242	M	26-35	20	A	;
11	1000005	P00014542	M	26-35	20	A	;
12	1000005	P00031342	M	26-35	20	A	;
13	1000005	P00145042	M	26-35	20	A	;
14	1000006	P00231342	F	51-55	9	A	;
15	1000006	P00190242	F	51-55	9	A	;
16	1000006	P0096642	F	51-55	9	A	;
17	1000006	P00058442	F	51-55	9	A	;
18	1000007	P00036842	M	36-45	1	B	;
19	1000008	P00249542	M	26-35	12	C	;

In [31]: `df.describe()`

Out[31]:

	User_ID	Occupation	Stay_In_Current_City_Years	Marital_Status	Product_Cate
count	5.500680e+05	550068.000000	550068.000000	550068.000000	550068.000000
mean	1.003029e+06	8.076707	1.858418	0.409653	5.409653
std	1.727592e+03	6.522660	1.289443	0.491770	3.931770
min	1.000001e+06	0.000000	0.000000	0.000000	1.000000
25%	1.001516e+06	2.000000	1.000000	0.000000	1.000000
50%	1.003077e+06	7.000000	2.000000	0.000000	5.000000
75%	1.004478e+06	14.000000	3.000000	1.000000	8.000000
max	1.006040e+06	20.000000	4.000000	1.000000	20.000000

In [34]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   User_ID                               550068 non-null  int64
1   Product_ID                           550068 non-null  object
2   Gender                               550068 non-null  object
3   Age                                   550068 non-null  object
4   Occupation                           550068 non-null  int64
5   City_Category                        550068 non-null  object
6   Stay_In_Current_City_Years          550068 non-null  int64
7   Marital_Status                      550068 non-null  int64
8   Product_Category                    550068 non-null  int64
9   Purchase                            550068 non-null  int64
dtypes: int64(6), object(4)
memory usage: 42.0+ MB
```

-columns like Age, Gender, City_Category are Categorical.

-column Purchase is Continuous.

-City_Category is of nominal type (no proper order between the categories).

-Age is of ordinal type (an order exists between the categories)

Converting Marital_Status to categorical

```
In [36]: # Get current data type of columns  
df.dtypes
```

```
Out[36]: User_ID          int64  
Product_ID        object  
Gender            object  
Age              object  
Occupation        int64  
City_Category     object  
Stay_In_Current_City_Years  int64  
Marital_Status    int64  
Product_Category  int64  
Purchase          int64  
dtype: object
```

```
In [37]: df['Marital_Status'] = pd.Categorical(df.Marital_Status)  
df.dtypes
```

```
Out[37]: User_ID          int64  
Product_ID        object  
Gender            object  
Age              object  
Occupation        int64  
City_Category     object  
Stay_In_Current_City_Years  int64  
Marital_Status    category  
Product_Category  int64  
Purchase          int64  
dtype: object
```

Converting Occupation to categorical

```
In [49]: df['Occupation'] = pd.Categorical(df.Occupation)  
df.dtypes
```

```
Out[49]: User_ID          int64  
Product_ID        object  
Gender            object  
Age              object  
Occupation        category  
City_Category     object  
Stay_In_Current_City_Years  int64  
Marital_Status    category  
Product_Category  int64  
Purchase          int64  
dtype: object
```

Shape of the data

```
In [42]: #shape of data  
shapeofdf = np.shape(df)  
print(shapeofdf)
```

```
(550068, 10)
```

```
In [43]: # The data set has 550068 rows and 10 columns
```

```
In [90]: max(df.Purchase)
```

```
Out[90]: 23961
```

Value counts and unique attributes


```
In [67]: cat_Count_Age = df["Age"].value_counts()  
print(cat_Count_Age)
```

```
Age  
26-35    219587  
36-45    110013  
18-25     99660  
46-50     45701  
51-55     38501  
55+       21504  
0-17      15102  
Name: count, dtype: int64
```

```
The count of purchase done by the age group 26-35 is highest
```

```
In [69]: x_bar=cat_Count_Age.index  
y_bar=cat_Count_Age  
plt.bar(x_bar,y_bar)
```

Out[69]: <BarContainer object of 7 artists>



```
In [51]: df['Age'].value_counts(normalize=True)
```

Out[51]: Age
26-35 0.399200
36-45 0.199999
18-25 0.181178
46-50 0.083082
51-55 0.069993
55+ 0.039093
0-17 0.027455
Name: proportion, dtype: float64

```
In [71]: cat_count_ms = df["Marital_Status"].value_counts()  
print(cat_count_ms)
```

Marital_Status
0 324731
1 225337
Name: count, dtype: int64

```
In [52]: df['Marital_Status'].value_counts(normalize=True)
```

Out[52]: Marital_Status
0 0.590347
1 0.409653
Name: proportion, dtype: float64

```
In [48]: df["City_Category"].value_counts()
```

```
Out[48]: City_Category
B      231173
C      171175
A      147720
Name: count, dtype: int64
```

```
In [53]: df['City_Category'].value_counts(normalize=True)
```

```
Out[53]: City_Category
B      0.420263
C      0.311189
A      0.268549
Name: proportion, dtype: float64
```

```
In [50]: df["Occupation"].value_counts()
```

```
Out[50]: Occupation
4      72308
0      69638
7      59133
1      47426
17     40043
20     33562
12     31179
14     27309
2      26588
16     25371
6      20355
3      17650
10     12930
5      12177
15     12165
11     11586
19      8461
13      7728
18      6622
9       6291
8       1546
Name: count, dtype: int64
```

```
In [54]: df['Occupation'].value_counts(normalize=True)
```

```
Out[54]: Occupation
4      0.131453
0      0.126599
7      0.107501
1      0.086218
17     0.072796
20     0.061014
12     0.056682
14     0.049647
2      0.048336
16     0.046123
6      0.037005
3      0.032087
10     0.023506
5      0.022137
15     0.022115
11     0.021063
19     0.015382
13     0.014049
18     0.012039
9      0.011437
8      0.002811
Name: proportion, dtype: float64
```

```
In [55]: df.Age.unique()
```

```
Out[55]: array(['0-17', '55+', '26-35', '46-50', '51-55', '36-45', '18-25'],
              dtype=object)
```

```
In [56]: df.head()
```


```
Out[56]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years
0	1000001	P00069042	F	0-17	10	A	2
1	1000001	P00248942	F	0-17	10	A	2
2	1000001	P00087842	F	0-17	10	A	2
3	1000001	P00085442	F	0-17	10	A	2
4	1000002	P00285442	M	55+	16	C	4

Visual Analysis - Univariate & Bivariate

```
In [58]: plt.figure(figsize=(10,8))  
sns.countplot(x='Age',hue='Gender',data=df)  
plt.ylabel('Purchase')
```

```
Out[58]: Text(0, 0.5, 'Purchase')
```



We can infer from above countplot:


The age group between 0 to 17 and above 55 there is less number of purchase at Walmart irrespective of the gender.

The number of purchase made by males is higher than females in all the age groups.

The highest number of purchase is done by the age group 26-35 males.

```
In [62]: plt.figure(figsize=(5,4))  
sns.countplot(x='Marital_Status',data=df)  
plt.ylabel('Purchase')
```


```
Out[62]: Text(0, 0.5, 'Purchase')
```



From the above countplot it can be seen that Purchase done by non married is more than the married people.

```
In [65]: plt.figure(figsize=(8,6))  
sns.countplot(x='Age',hue='Marital_Status',data=df)  
plt.ylabel('Purchase')
```

```
Out[65]: Text(0, 0.5, 'Purchase')
```




It can be inferred from the above countplot that highest number of purchase done by non married people in the age group 26-35. And the least by the non married of age of age group above 55 and below 17.

Histplot

```
In [86]: sns.histplot(df['Purchase'],bins = 10)
```

```
Out[86]: <Axes: xlabel='Purchase', ylabel='Count'>
```




From the histplot it can be seen that the number of purchase made in the range of 5000 to 10000 is highest

KDE Plot

```
In [94]: sns.kdeplot(df['Purchase'])
```


```
Out[94]: <Axes: xlabel='Purchase', ylabel='Density'>
```



BoxPlot

```
In [82]: plt.figure(figsize=(10,5))  
sns.boxplot(y = df["Purchase"])  
plt.yticks(fontsize=20)  
plt.ylabel('Purchase', fontsize=20)  
plt.title('Purchase', fontsize=20)
```


```
Out[82]: Text(0.5, 1.0, 'Purchase')
```



Minimum, excluding outliers: 0
Maximum, excluding outliers: 20,000 above
25th Quantile: 5000
Median: around 7000
75th Quantile: 12500


There are few outliers

```
In [88]: sns.boxplot(x='Age', y='Purchase', data=df)
plt.xticks(rotation=90, fontsize=12)
plt.yticks(fontsize=12)
plt.title('Purchase Age wise', fontsize=15)
plt.show()
```



The overall Purchase of all the age category , has similar spread.(5000 to 12000)
All the age group have many outliers.
the median is almost the same for all the age group.

```
In [92]: sns.boxplot(x='Gender', y='Purchase', data=df)
plt.xticks(rotation=90, fontsize=12)
plt.yticks(fontsize=12)
plt.title('Purchase Gender wise', fontsize=15)
plt.show()
```



The overall Purchase of both the gender , has similar spread.(5000 to 12000)

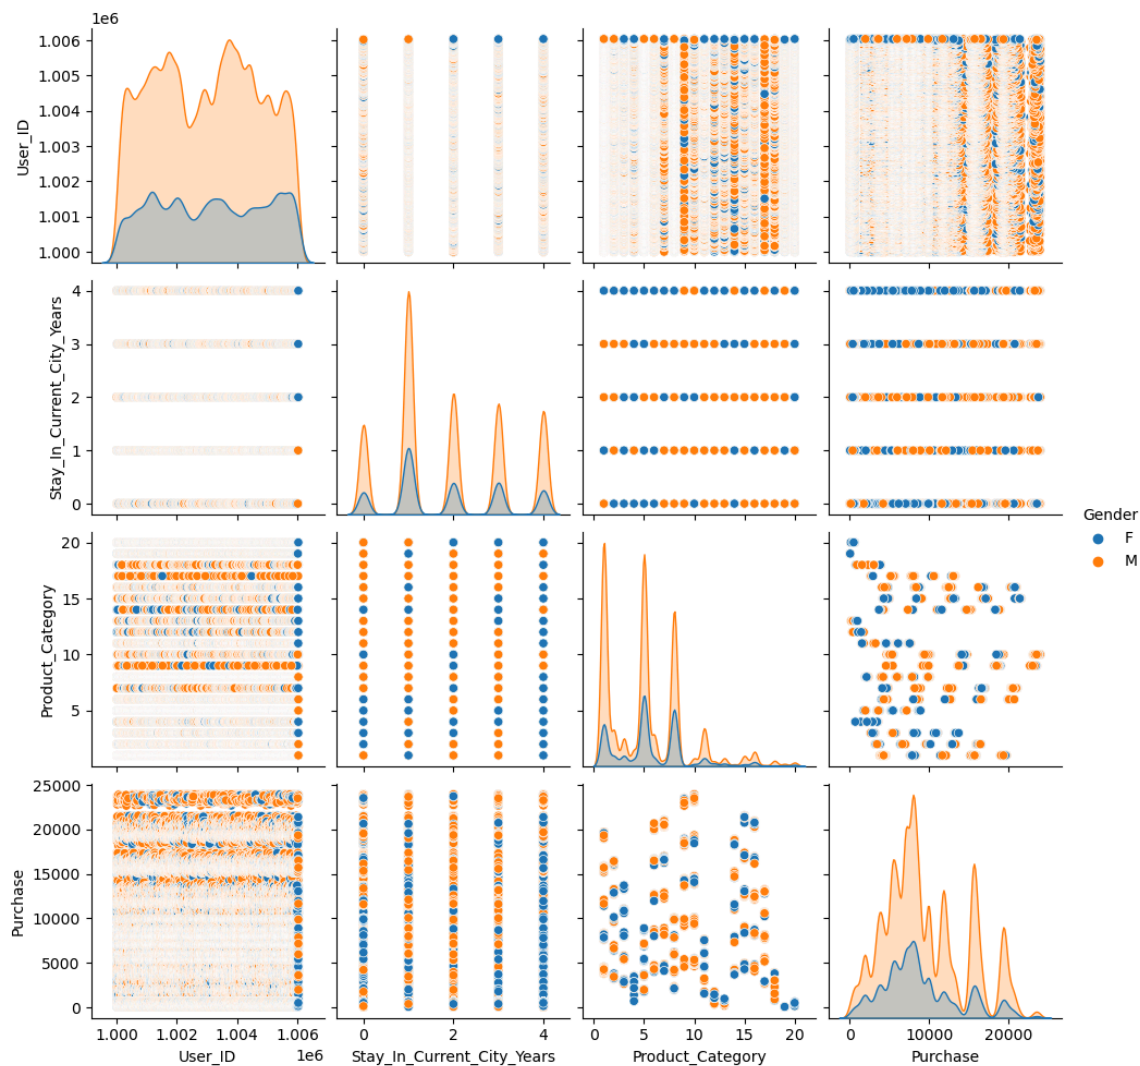
Both the gender have many outliers.

the median is almost the same for both the gender.

Pairplot

```
In [99]: sns.pairplot(data=df, hue='Gender')
plt.show()
```

C:\Users\Hp\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)




Correlation Matrix

```
In [100]: #to find the level of correlation b/w different attributes (variables)
```



```
In [102]: sns.scatterplot(x= 'Purchase', y= 'Age', data = df)
plt.show()
```



Heatmap

Missing Value & Outlier Detection

```
In [104]: df.isnull().values.any()
```

Out[104]: False

```
In [105]: df.isnull().sum()
```


```
Out[105]: User_ID          0
Product_ID         0
Gender             0
Age               0
Occupation         0
City_Category      0
Stay_In_Current_City_Years  0
Marital_Status     0
Product_Category   0
Purchase           0
dtype: int64
```

There are no missing values

In [108]: *#Outliers can be detected using boxplot*

```
In [106]: plt.figure(figsize=(10,5))
sns.boxplot(y = df["Purchase"])
plt.yticks(fontsize=20)
plt.ylabel('Purchase', fontsize=20)
plt.title('Purchase', fontsize=20)
```

Out[106]: Text(0.5, 1.0, 'Purchase')



There are many outliers and they can be treated using various techniques

1. Minimum Score It is the lowest value 5000.
2. Lower Quartile 25% of values fall below the lower quartile value. It is also known as the first quartile.
3. Median Median marks the mid-point of the data. It is shown by the line that divides the box into two parts. Half the scores are greater than or equal to this value and half are less. It is sometimes known as the second quartile.
4. Upper Quartile 75% of values fall below the upper quartile value. It is also known as the third quartile. Maximum Score It is the highest value, excluding outliers. It is shown at the end of upper whisker. Whiskers The upper and lower whiskers represent values outside the middle 50%. That is, the lower 25% of values and the upper 25% of values. Interquartile Range (or IQR) This is the box plot showing the middle 50% of scores. It is the range between the 25th and 75th percentile.

Do some data exploration steps like:

-Tracking the amount spent per transaction of all the 50 million female customers, and all the 50 million male customers, calculate the average, and conclude the results.

-Inference after computing the average female and male expenses.

-Use the sample average to find out an interval within which the population average will lie. Using the sample of female customers you will calculate the interval within which the average spending of 50 million male and female customers may lie.

```
In [112]: grouped = df.groupby('Gender')['Purchase'].sum()  
print(grouped)
```

```
Gender  
F    1186232642  
M    3909580100  
Name: Purchase, dtype: int64
```

```
In [113]: df["Purchase"].sum()
```

```
Out[113]: 5095812742
```

```
In [117]: total_purchase_male = 3909580100  
total_purchase_female = 1186232642  
  
total_purchase = 5095812742  
  
# Calculate percent purchase for female customers  
percentage_female = total_purchase_female / total_purchase  
  
# Calculate percent amount purchase for male customers  
percentage_male = total_purchase_male / total_purchase
```

```
In [118]: print(percentage_female)
```


```
0.2327857600070344
```

```
In [119]: print(percentage_male)
```

```
0.7672142399929656
```

```
In [121]: plt.figure(figsize=(6,4))  
sns.countplot(x='Gender',data=df)  
plt.ylabel('Purchase')
```

```
Out[121]: Text(0, 0.5, 'Purchase')
```



```
In [122]: count1 = df.groupby('Gender')['Purchase'].count()  
print(count1)
```

```
Gender  
F    135809  
M    414259  
Name: Purchase, dtype: int64
```

```
In [123]: num_transactions_female = 135809  
num_transactions_male = 414259
```

```
In [124]: #Calculate the average amount spent per transaction  
average_female = total_purchase_female / num_transactions_female  
average_male = total_purchase_male / num_transactions_male
```

```
In [127]: print(average_female,average_male)
```

```
8734.565765155476 9437.526040472265
```

women are not spending more money per transaction than men.

Since the number of transactions by women and also the amount spent on the purchase is less compared to men

Use the sample average to find out an interval within which the population average will lie.

Using the sample of female customers you will calculate the interval within which the average spending of 50 million male and female customers may lie.

To calculate the confidence interval for the population average spending of female customers and male customers based on the sample averages, we can use the t-distribution since we are dealing with sample averages and want to estimate the population mean.

$z_1 = \text{norm.ppf}(0.025)$ as we have 2.5% data till z_1
Similarly to calculate z_2 will use

$z_2 = \text{norm.ppf}(1-0.025)$ as we have 2.5% data remaining after z_2
We can also represent the z_1 points as 2.5th percentile and z_2 as 97.5th percentile ($1-0.025$)

```
In [194]: # z1 will be
z1 = norm.ppf(0.025)
z1
```

```
Out[194]: -1.9599639845400545
```

```
In [195]: # z2 will be
z2 = norm.ppf(1 - 0.025) # we can also use norm.ppf(0.975)
z2
```

```
Out[195]: 1.959963984540054
```

$$Z = \frac{\bar{X} - \mu}{\sigma/\sqrt{n}}$$

Where,

\bar{X} = sample mean μ - population mean σ/\sqrt{n} = standard error Here in the above example our sample S, the population mean will lie between z scores i.e. z_1 and z_2 which is -1.96 and 1.96.

Between these two points, we will have our 95% confidence interval.

```
In [ ]: ## Female CI
```

```
In [137]: import math
```

```
In [229]: # Calculate the mean
mean_female = 1186232642 / 135809

# Calculate the standard deviation
standard_deviation_female = (math.sqrt(mean_female / 135809))
```

```
In [230]: print(mean_female)
```

```
8734.565765155476
```

```
In [231]: print(standard_deviation_female)
```

```
0.25360416620895176
```

```
In [232]: # std deviation  sigma/sqrt(n)
std_error_female = .25/np.sqrt(135809)
std_error_female
```

```
Out[232]: 0.0006783842134393103
```

to get the data points for z1 which is on left side and z2 which is on right side,

$X1 = \mu + Z1 * \sigma / \sqrt{n}$ and,

$X2 = \mu + Z2 * \sigma / \sqrt{n}$

```
In [233]: x1 = 8734.56 + z1 * std_error_female
x1
```

```
Out[233]: 8734.558670391374
```

```
In [234]: x2 = 8734.56 + z2 * std_error_female
x2
```

```
Out[234]: 8734.561329608625
```

So the range of 95% confidence interval --> [8734.42, 8734.69]

```
In [235]: ## Using norm.interval()
```

```
In [236]: norm.interval(0.95, loc=8734.56, scale=std_error)
```

```
Out[236]: (8734.559558449768, 8734.560441550231)
```

```
In [237]: ## male CI
```

```
In [238]: # Calculate the mean
mean_male = 3909580100 / 414259
```

```
In [239]: # Calculate the standard deviation
standard_deviation_male = (math.sqrt(mean_female / 414259))
```

```
In [240]: print(mean_male)
```

```
9437.526040472265
```

```
In [241]: print(standard_deviation_male)
```

```
0.14520604054667763
```

```
In [203]: # std deviation sigma/sqrt(n)
std_error_male = 0.145/np.sqrt(414259)
std_error_male
```

Out[203]: 0.00022528487013342337

to get the data points for z1 which is on left side and z2 which is on right side,

$$X1 = \mu + Z1 * \sigma / \sqrt{n} \text{ and,}$$

$$X2 = \mu + Z2 * \sigma / \sqrt{n}$$

```
In [204]: xm1 = 9437.5 + z1 * std_error_male
xm1
```

Out[204]: 9437.499558449768

```
In [205]: xm2 = 9437.5 + z2 * std_error_male
xm2
```

Out[205]: 9437.500441550232

In []: So the range of 95% confidence interval --> [9437.49, 9437.500]

In []: ## Using norm.interval()

```
In [170]: norm.interval(0.95, loc=9437.5, scale=std_error)
```

Out[170]: (9437.499558449768, 9437.500441550232)

The confidence interval (CI) gives us a range of values within which we are confident that the average spending per transaction lies. It provides measure of the uncertainty.

Based on the confidence intervals for the average spending of both female and male customers, we can make some inferences based on these intervals. Here are a few points of inference:

Female customers: 8734.42, 8734.69

Male customers: 9437.49, 9437.50

Since there is no overlapping in the confidence intervals for the two groups, it suggests that there is a significant difference in the average spending between the two groups. We can conclude with 95% confidence that there is no statistically significant difference in the average spending per transaction between female and male customers. It might be due to the variability within the samples.

It's also essential to consider other factors that might influence spending behavior, such as age, location, type of products purchased, promotions, etc.

In []:

Same way fro 99% Confident Interval

```
In [172]: z1=norm.ppf(0.005) #as we have 0.5% data till z1
```

```
In [173]: #Similarly to calculate z2 will use z2=norm.ppf(1-0.005)
```

as we have 0.5% data remaining after z2 We can also represent the z1 points as 0.5th percentile and z2 as 99.5th percentile (1-0.005)

```
In [188]: # z1 will be
zz1 = norm.ppf(0.005)
zz1
```

```
Out[188]: -2.575829303548901
```

```
In [189]: # z2 will be
zz2 = norm.ppf(1 - 0.005) # we can also use norm.ppf(0.975)
zz2
```

```
Out[189]: 2.5758293035489004
```

$$Z = \frac{\bar{X} - \mu}{\sigma/\sqrt{n}}$$

Where,

\bar{X} = sample mean

μ - population mean

σ/\sqrt{n} = standard error

Here in the above example our sample S, the population mean will lie between z scores i.e. z1 and z2 which is -2.57 and 2.57.

Between these two points, we will have our 99% confidence interval.

```
In [177]: ## Female CI
```

```
import math
```

```
In [178]: # Calculate the mean
mean_female = 1186232642 / 135809
```

```
In [179]: # Calculate the standard deviation
standard_deviation_female = (math.sqrt(mean_female / 135809))*100
```

```
In [180]: print(mean_female)
```

```
8734.565765155476
```

```
In [181]: print(standard_deviation_female)
```

```
25.360416620895176
```



```
In [209]: # std deviation sigma/sqrt(n)
std_error_female = 25.3/np.sqrt(135809)
std_error_female
```

Out[209]: 0.0686524824000582

to get the data points for z1 which is on left side and z2 which is on right side,

$$X1 = \mu + Z1 * \sigma / \sqrt{n} \quad \text{and,}$$

$$X2 = \mu + Z2 * \sigma / \sqrt{n}$$

```
In [210]: x1 = 8734.56 + zz1 * std_error_female
x1
```

Out[210]: 8734.383162924072

```
In [211]: x2 = 8734.56 + zz2 * std_error_female
x2
```

Out[211]: 8734.736837075927

In []: So the range of 99% confidence interval --> [8734.38, 8734.73]

```
In [226]: ## Using norm.interval()

norm.interval(0.99, loc=8734.56, scale=std_error)
```

Out[226]: (8734.559419704628, 8734.56058029537)

male CI

```
In [213]: # Calculate the mean
mean_male = 3909580100 / 414259
```

```
In [214]: # Calculate the standard deviation
standard_deviation_male = (math.sqrt(mean_female / 414259))
```

```
In [215]: print(mean_male)

9437.526040472265
```

```
In [216]: print(standard_deviation_male)

0.14520604054667763
```

```
In [217]: # std deviation sigma/sqrt(n)
std_error_male = 0.145/np.sqrt(414259)
std_error_male
```

Out[217]: 0.00022528487013342337

to get the data points for z1 which is on left side and z2 which is on right side,

$X1 = \mu + Z1 * \sigma / \sqrt{n}$ and,

$X2 = \mu + Z2 * \sigma / \sqrt{n}$

```
In [222]: xm1 = 9437.5 + zz1 * std_error_male
xm1
```

```
Out[222]: 9437.499419704629
```

```
In [223]: xm2 = 9437.5 + zz2 * std_error_male
xm2
```

```
Out[223]: 9437.500580295371
```

```
In [ ]: So the range of 95% confidence interval --> [9437.49, 9437.50]
```

```
In [224]: ## Using norm.interval()

norm.interval(0.99, loc=9437.5, scale=std_error)
```

```
Out[224]: (9437.499419704629, 9437.500580295371)
```

The confidence interval (CI) gives us a range of values within which we are confident that the average spending per transaction lies. It provides measure of the uncertainty.

Based on the confidence intervals for the average spending of both female and male customers, we can make some inferences based on these intervals. Here are a few points of inference:

```
In [ ]: Female customers: 8734.38, 8734.73
```

```
Male customers: 9437.49, 9437.50
```

Since there is no overlapping in the confidence intervals for the two groups, it suggests that there is no significant difference in the average spending between the two groups. We can conclude with 95% confidence that there is no statistically significant difference in the average spending per transaction between female and male customers. It might be due to the variability within the samples.

It's also essential to consider other factors that might influence spending behavior, such as age, location, type of products purchased, promotions, etc.

There is no difference when confidence interval is changed from 95 to 99.

the results and check if the confidence intervals of married and non married spends are overlapping or not overlapping. How can Walmart leverage this conclusion to make changes or improvements?

In [227]: `df.head()`

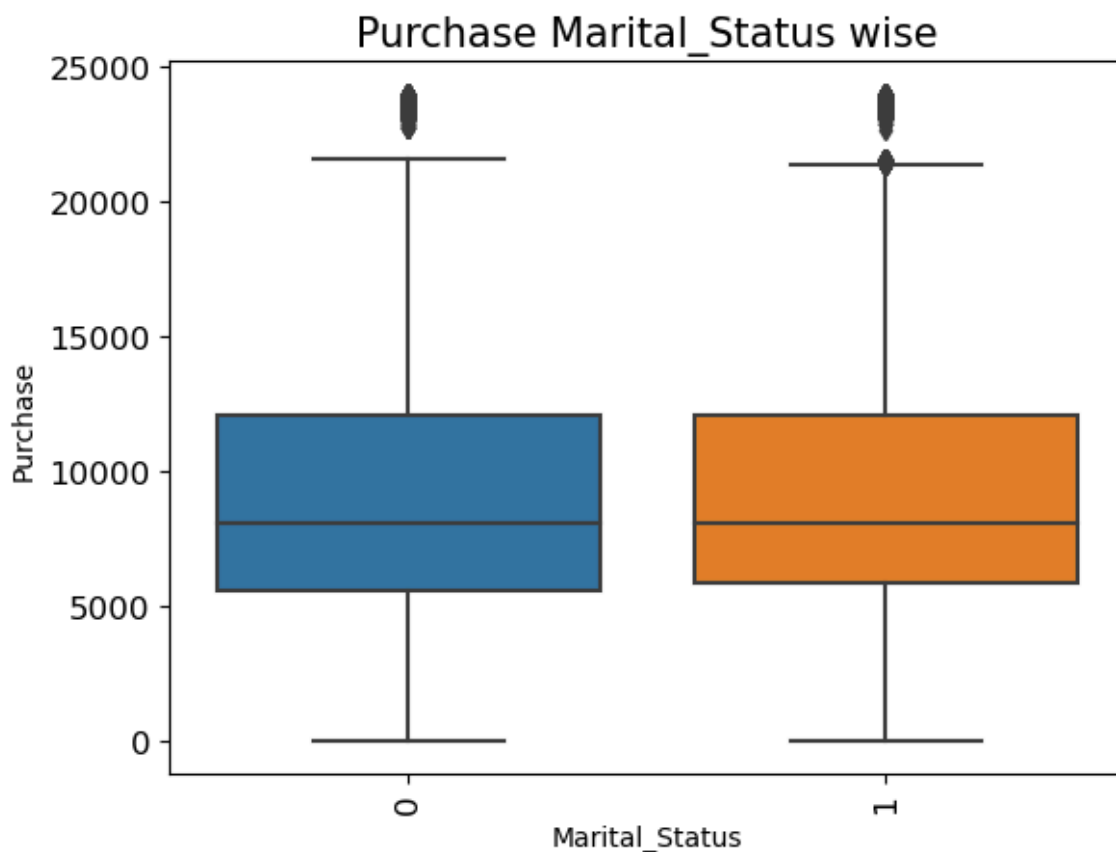
Out[227]:

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years
0	1000001	P00069042	F	0-17	10	A	2
1	1000001	P00248942	F	0-17	10	A	2
2	1000001	P00087842	F	0-17	10	A	2
3	1000001	P00085442	F	0-17	10	A	2
4	1000002	P00285442	M	55+	16	C	4

In [228]: `count_MS = df.groupby('Marital_Status')['Purchase'].count()
print(count_MS)`

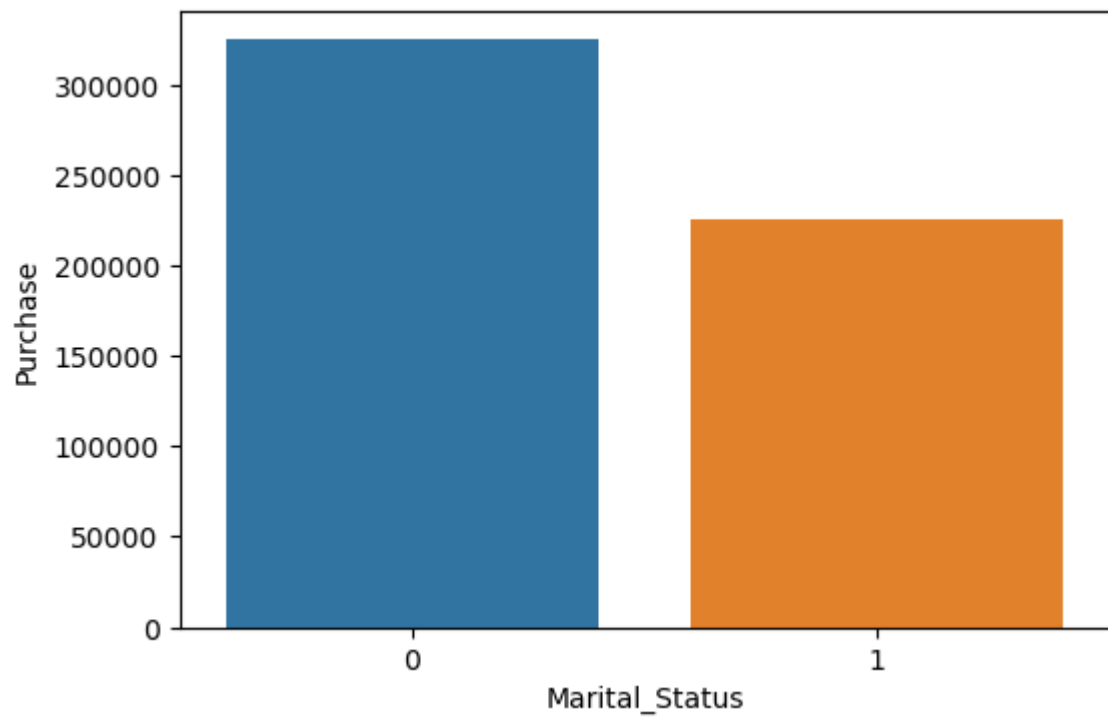
```
Marital_Status
0      324731
1      225337
Name: Purchase, dtype: int64
```

In [243]: `sns.boxplot(x='Marital_Status', y='Purchase', data=df)
plt.xticks(rotation=90, fontsize=12)
plt.yticks(fontsize=12)
plt.title('Purchase Marital_Status wise', fontsize=15)
plt.show()`



```
In [244]: plt.figure(figsize=(6,4))  
sns.countplot(x='Marital_Status',data=df)  
plt.ylabel('Purchase')
```

Out[244]: Text(0, 0.5, 'Purchase')



In []: